

Anh Dao

anhdd2000@gmail.com | +358 41 325 1677 | anh.d.dao@aalto.fi | Portfolio

Education

M.Sc. in Computer Science

Aalto University, Finland

August 2024 – Current

GPA: 4.25 / 5.00

B.A. in Computer Science

Grinnell College, United States

August 2019 – May 2023

GPA: 3.76 / 4.00

Relevant Coursework: Information Security, Network Security, Cloud Software and Development, Deep Learning, Machine Learning, Scalable Systems, Biomedical Machine Learning, Linear Optimization, Web Software Development, Probability and Statistics

Experience

Nokia Solutions & Networks Oy – Finland

August 2025 – Current

R&D Trainee (*Python/React/GCP*)

- Developing multi-agentic systems and document processing pipelines utilizing Python, React, GCP, and Vertex AI to create data corpora for RAG systems, targeting high-accuracy retrieval for technical documentation
- Building LLM-powered internal tools that automated data processing workflows, reducing manual effort by approximately 40% for the 4LS team operations
- Engineered scalable RAG agents capable of processing and indexing department-wide knowledge bases, streamlining information access for internal stakeholders

Aalto University – Finland

August 2024 – Current

Teaching Assistant (*Python/Docker/Kubernetes*) – Computer Science Department

- Provided technical support and grading for 5 Master-level security courses, directly assisting over 300 students and achieving a consistent 4.5/5.0 feedback rating
- Deployed online courses through the A+ Learning Management System, utilizing Docker/Kubernetes to serve containerized lab environments with 99.9% uptime during peak submission periods
- Developed and optimized Python-based autograders, reducing grading turnaround time by 60% and ensuring immediate feedback for students
- Modified CS-E4300 Network Security exercises based on feedback, resolving 150+ technical tickets and improving clarity on complex network configuration tasks

Microsec Zrt. – Budapest, Hungary

June 2022 – August 2022

Software Engineer Trainee (*Go*)

- Developed an update to the company's PKI certificate reading linter program to properly handle edge cases of subordinate PKI certificates following RFC5280 standards
- Implemented parsing logic for X.509 certificate extensions and validation rules

Grinnell College – United States

August 2022 – May 2023

Research Assistant (*R/Python*) – Data Analysis and Social Inquiry Lab

- Built multivariate statistical models in R using data gathered from games on Stat2Labs, used by over 1000 students in undergraduate statistics courses across 5 higher education institutions
- Helped on average over 150 students and faculty weekly with computing and data-related projects

Grinnell College – United States

August 2021 – December 2022

Peer Educator (*Java/C*) – Computer Science Department

- Hosted weekly mentor sessions for 100+ students across multiple CS courses, providing coding challenges and interactive Q&A sessions on course materials

Course Projects

Complex Event Processing with Proper Load Shedding (Python) September 2025 – December 2025

- Implemented state-based load shedding strategy for the OpenCEP framework to handle bursty workloads while maintaining low latency and acceptable recall
- Developed adaptive load shedding using weighted scoring functions considering importance, priority, time criticality, station importance, and chain potential
- Designed state management for partial matches using hash table implementations achieving O(1) retrieval time complexity
- Implemented hot-path detection patterns using Kleene closure operators on CitiBike dataset (~1.8M trips, 10,850 unique bikes)
- Achieved 90-100% recall rates while maintaining latency limits through intelligent partial match pruning
- Developed feedback controller for dynamic drop fraction adjustment based on p90 latency and CPU/memory pressure
- *Role: Semantic/hybrid load-shedding and state-management implementation. Collaborators: Quoc-Huy Trinh, Pablo Rubio*
- *Code: github.com/huyquoctrinh/ScalableSys-Assignment1*

Enhance Graph Retrieval-Augmented Generation by Caching Method (Python/Kuzu) September 2025 – December 2025

- Enhanced GraphRAG pipeline by integrating LRU caching mechanism and memory-efficient caching variant combining context pruning and on-the-fly compression
- Implemented Text2Cypher component using dynamic few-shot prompting with TF-IDF-based exemplar selection for semantic similarity matching
- Developed iterative query refinement with database feedback loop for error correction and deterministic post-processing for schema validation
- Achieved highest throughput and lowest latency at scale with memory-efficient LRU caching, demonstrating 60-90% memory savings for large context-heavy entries
- Utilized Kuzu graph database for heterogeneous graph storage with nodes (Laureates, Scholars, Prizes, Institutions) and edges (AFFILIATED_WITH, AWARDED)
- *Role: Few-shot prompting implementation and dataset creation. Collaborators: Quoc-Huy Trinh, Pablo Rubio*
- *Code: github.com/huyquoctrinh/ScalableSys-Proj2*

LLM Cost Analysis and Optimization (Python/GCP/Vertex AI) September 2025 – December 2025

- Analyzed LLM inference costs and optimization strategies for enterprise deployments
- Developed cost decomposition framework for understanding token-level pricing and latency tradeoffs
- Implemented benchmarking suite comparing various model configurations on Vertex AI
- *Code: github.com/daoa0601/cs-e4660-advanced-topics-software*

ICS Network Security in the Rise of IIoT Technology (Seminar Paper) January 2025 – April 2025

- Reviewed and analyzed the Purdue Enterprise Reference Architecture (PERA) model for Industrial Control Systems network security and zone-based segmentation
- Examined Zero-trust Architecture (ZTA) principles and implementation challenges in Operational

Technology (OT) environments

- Analyzed network segmentation methods including identity-based segmentation with blockchain, micro-segmentation for east-west traffic control, and Software-Defined Networking (SDN)
- Investigated Manufacturing Usage Descriptions (MUD) for IIoT device security and automated intrusion detection
- Reviewed IT/OT convergence trends and the applicability of machine learning for dynamic access policy enforcement
- *Tutor: Mohit Sethi*

Encrypted File Transfer Protocol (Python/PyCryptodome)

March 2022 – April 2022

- Implemented a secure file messaging protocol using PyCryptodome with Symmetric Key Cryptography following AES-GCM encryption mode
- Designed protocol to ensure message security even against server-side attacks through authenticated encryption

Sudoku Solver (C/CUDA)

March 2021

- Developed an automated Sudoku solving program using CUDA to perform parallel computing on over 100 Sudoku boards simultaneously
- Implemented constraint propagation and backtracking algorithms optimized for GPU execution

Skills

Languages: Python, C/C++, Java, Go, JavaScript/TypeScript, R, MATLAB, HTML/CSS, SQL, Cypher

Frameworks & Libraries: PyTorch, React, Node.js, Express.js, DSPy, RDKit, scikit-learn, pandas, NumPy

Tools & Platforms: Git, Docker, Kubernetes, Google Cloud Platform, Vertex AI, Wireshark, CUDA, PyCryptodome

Databases: Kuzu (Graph DB), Neo4j, PostgreSQL

Concepts: Machine Learning, Deep Learning, Graph Neural Networks, Diffusion Models, RAG Systems, Complex Event Processing, Network Security, Zero-trust Architecture, PKI/X.509, Cryptography